



(<http://vision.stanford.edu/>) (<http://stanford.edu/>)



CS231n: Convolutional Neural Networks for Visual Recognition

(<http://cs231n.stanford.edu/>)

Course Project Reports: Winter 2016 [Tweet](#)

(<https://twitter.com/share>)

Showing 130 out of 200 projects (70 were requested to remain private):

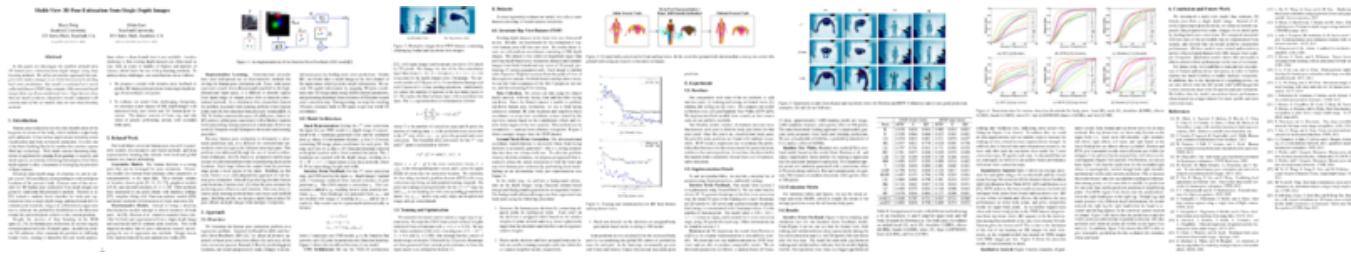
DeepPaint: A Tool For Image Inpainting
 (reports2016/211_Report.pdf)
 Suleman Kazi, Kushagr Gupta, Terry Kong



Attentional Scene Classification with Human Eye Movements (reports2016/004_Report.pdf)
 Alex Kuefler



Multi-View 3D Pose Estimation from Single Depth Images
 (reports2016/412_Report.pdf)
 Boya Peng, Zelun Luo



Supporting humanitarian mapping with deep learning (reports2016/409_Report.pdf)

Lars Roemheld

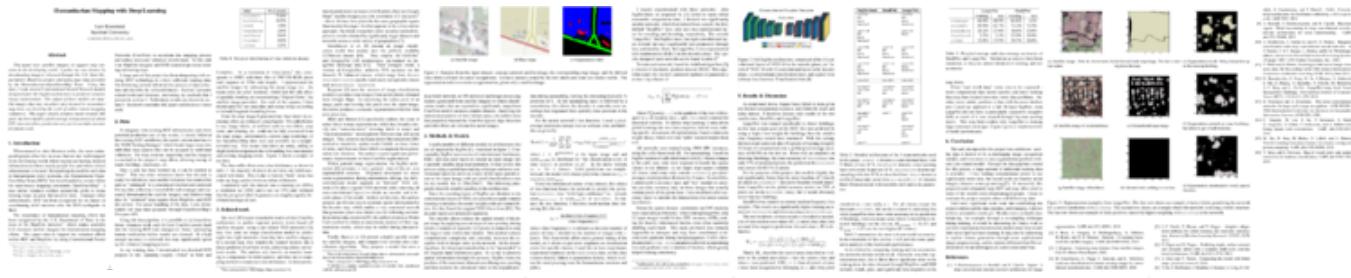
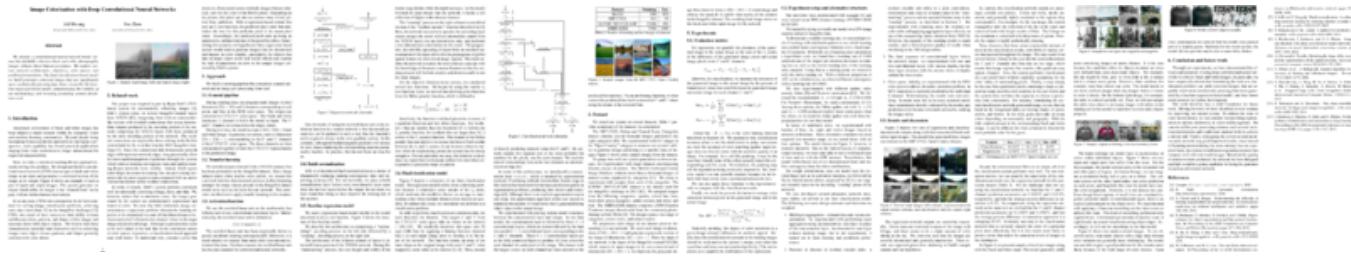


Image Colorization with Deep Convolutional Neural Networks (reports2016/219_Report.pdf)

Jeff Hwang, You Zhou



The Game Imitation: A Portable Deep Learning Model for Modern Gaming AI (reports2016/113_Report.pdf)

Zhao Chen, Darvin Yi



Stochastic Video Prediction with Deep Conditional Generative Models (reports2016/215_Report.pdf)

Rui Shu



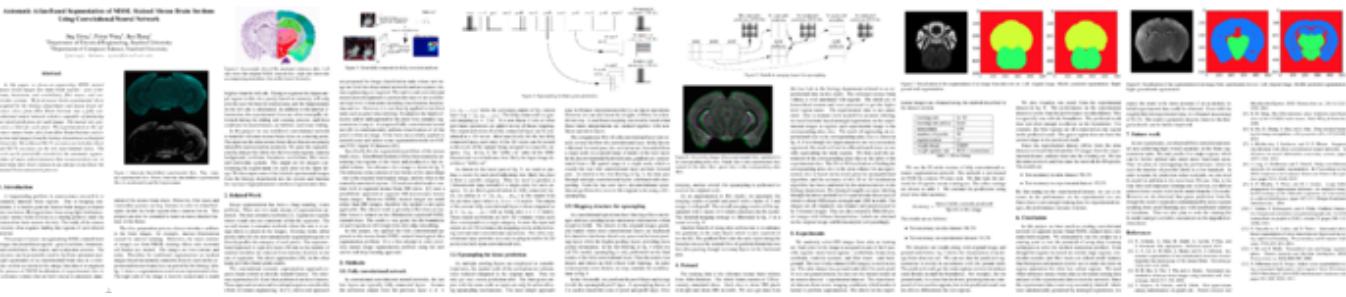
Detecting Diabetic Retinopathy (reports2016/309_Report.pdf)

Tanner Gilligan, Marco Alban



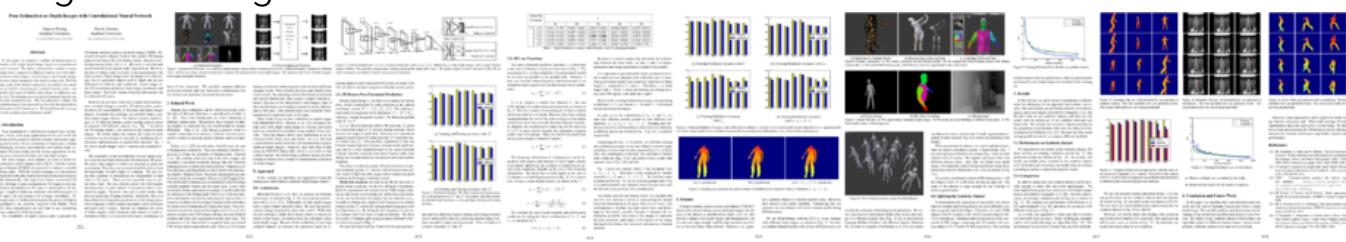
Automatic atlas-based segmentation of NISSL stained mouse brain sections using convolutional neural network (reports2016/318_Report.pdf)

Jing Xiong, Feiran Wang, Jian Zhang



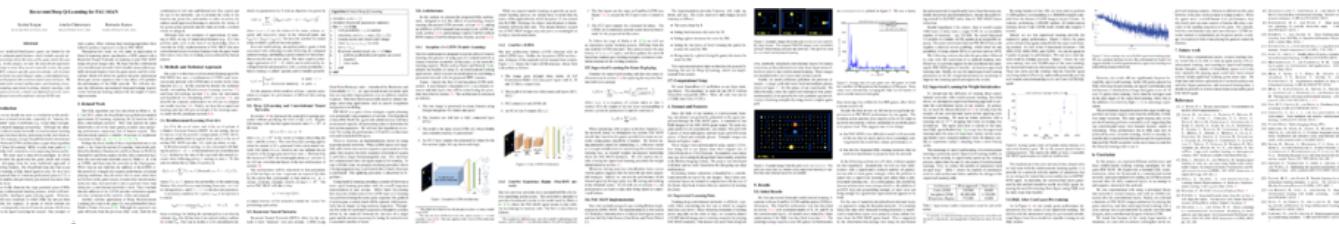
Pose Estimation on Depth Images with Convolutional Neural Network (reports2016/426_Report.pdf)

Jingwei Huang



Deep Q-Learning for PAC-MAN (reports2016/106_Report.pdf)

Bernardo Ramos, Kushal Ranjan, Amelia Christensen

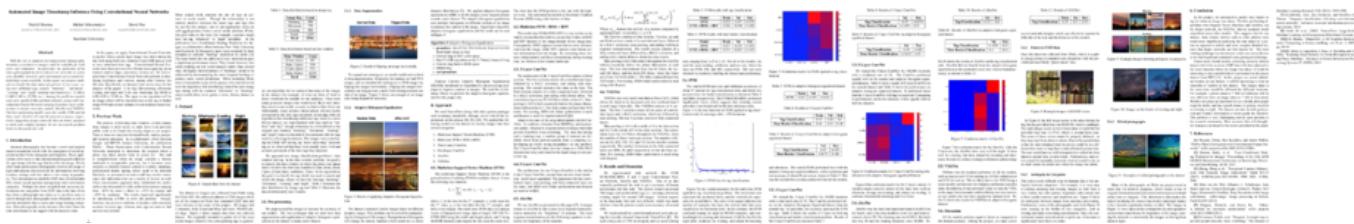


Recognizing Handwritten Characters (reports2016/269_Report.pdf)

Lisa Yan

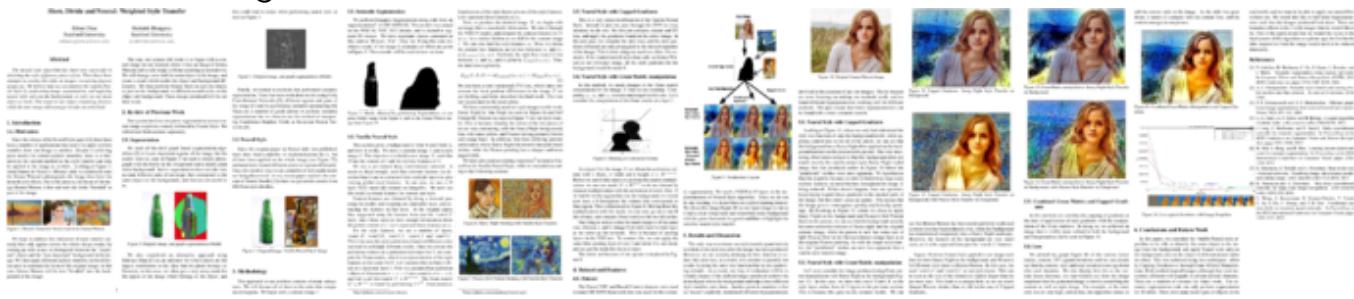
Automated Image Timestamp Inference using Convolutional Neural Networks (reports2016/267_Report.pdf)

Prafull Sharma, Michel Shoemaker, David Pan



Show, Divide and Neural: Weighted Neural Style Transfer (reports2016/208_Report.pdf)

Rishabh Bhargava, Ethan Chan



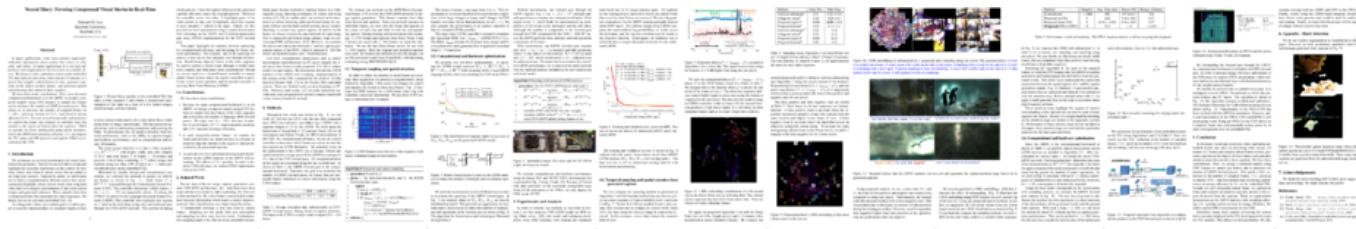
Neural Fill (reports2016/209_Report.pdf)

Russell Kaplan, Christopher Sauer, Alexander Lin



Neural Diary: Forming Compressed Visual Stories in Real-Time (reports2016/213_Report.pdf)

Edward Lee



Model-Based Reinforcement Learning for Playing Atari Games (reports2016/116_Report.pdf)

Irving Hsu, Justin Fu



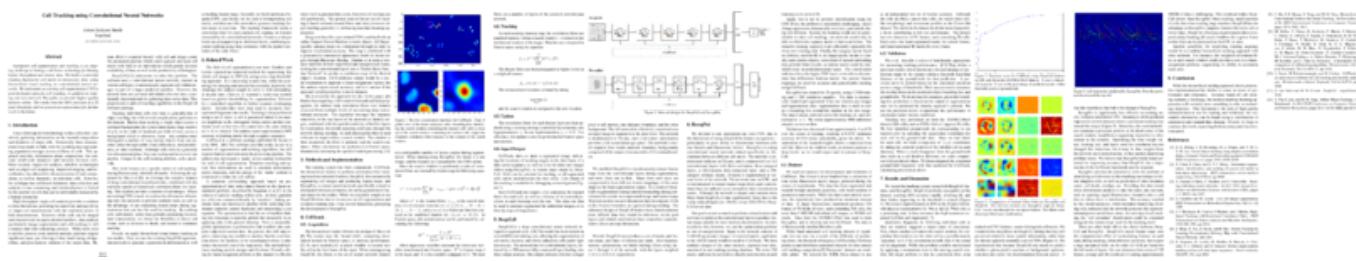
Deep Reinforcement Learning for Simulated Autonomous Vehicle Control (reports2016/112_Report.pdf)

Raphael Palefsky-Smith, April Yu, Rishi Bedi



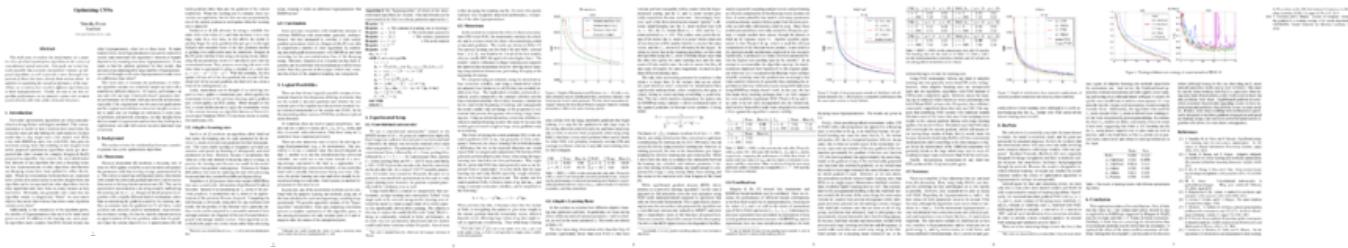
Cell Nucleus Tracking using Convolutional Neural Networks (reports2016/326_Report.pdf)

Anton Jackson-Smith



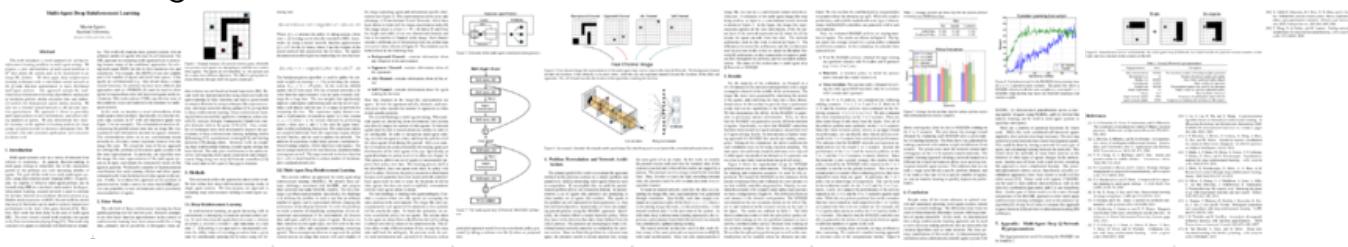
Optimizing Convolutional Networks (reports2016/108_Report.pdf)

Timothy Dozat



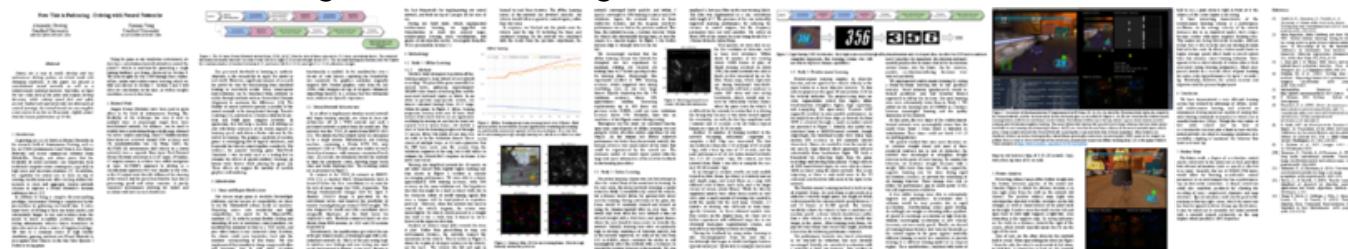
Multi-Agent Deep Reinforcement Learning (reports2016/122_Report.pdf)

Maxim Egorov



Now This is Pod Racing (reports2016/100_Report.pdf)

Alexander Dewing, Xiaonan Tong



Deep Learning of Spatial and Temporal Features for Automotive Prediction (reports2016/102_Report.pdf)

Tim Wheeler, Jeremy Morton



Complexity Beyond the Trigram: Identifying Sign Languages from Video (reports2016/207_Report.pdf)

Pamela Toman



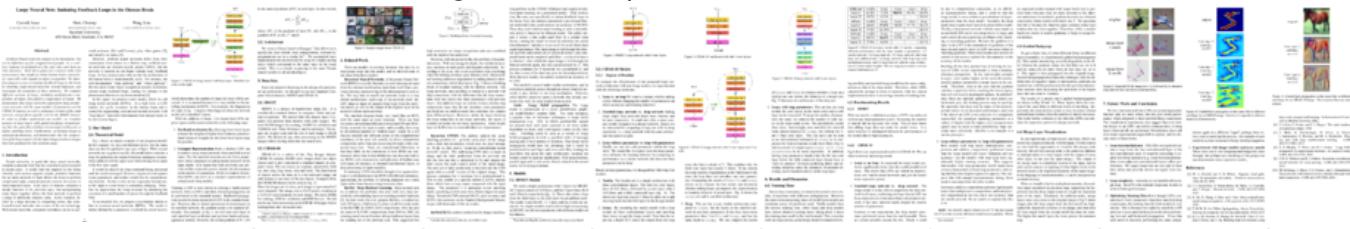
Signature Verification with Convolutional Neural Networks (reports2016/276_Report.pdf)

Gabriel Alvarez, Blue Sheffer, Morgan Bryant



Loopy Neural Nets: Imitating Feedback Loops in the Human Brain (reports2016/110_Report.pdf)

Isaac Caswell, Lisa Wang, Chuanqi Shen



DeepFace: Face Generation using Deep Learning (reports2016/006_Report.pdf)

Hardie Cate, Fahim Dalvi, Zeshan Hussain



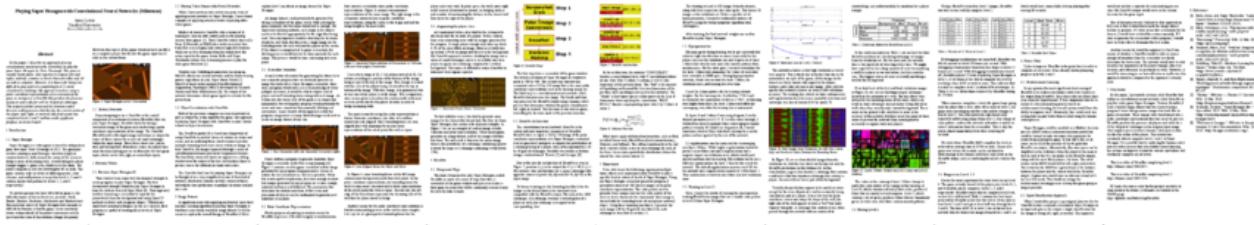
Pruning of Winograd and FFT Convolution Algorithm (reports2016/117_Report.pdf)

Xingyu Liu, Yatish Turakhia



Playing Super Hexagon with Convolutional Neural Networks (reports2016/115_Report.pdf)

Jason Lewis



Heirarchical Deep Convolutional Neural Networks: An Insect Classification Application (reports2016/283_Report.pdf)

Jeffrey Glick, Katarina Miller



Convolutional Neural Networks for Left Ventricle Volume Estimation (reports2016/305_Report.pdf)

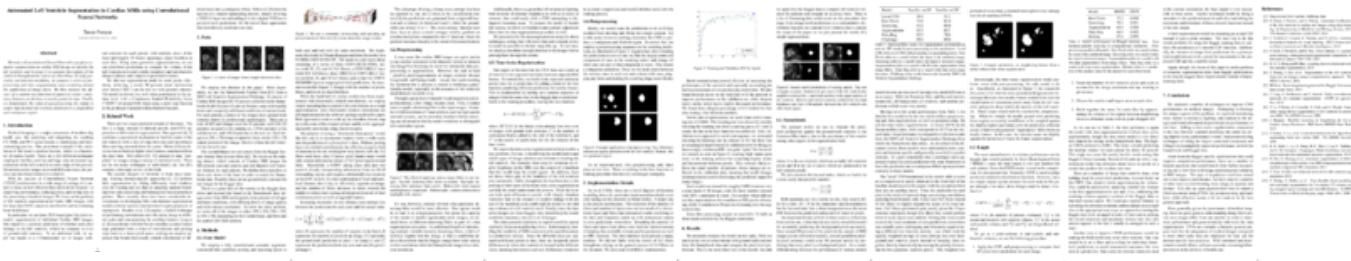
Cheryl Danner, Carol Hsin



Automated Left Ventricle Segmentation in Cardiac MRIs with Convolutional Neural Networks

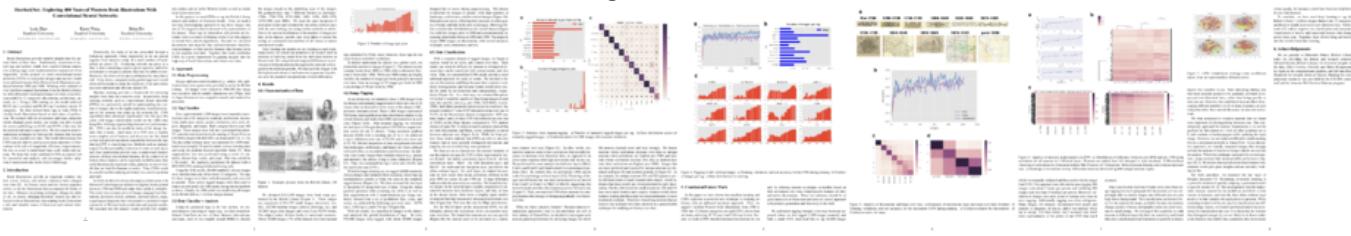
(reports2016/317_Report.pdf)

Taman Narayan



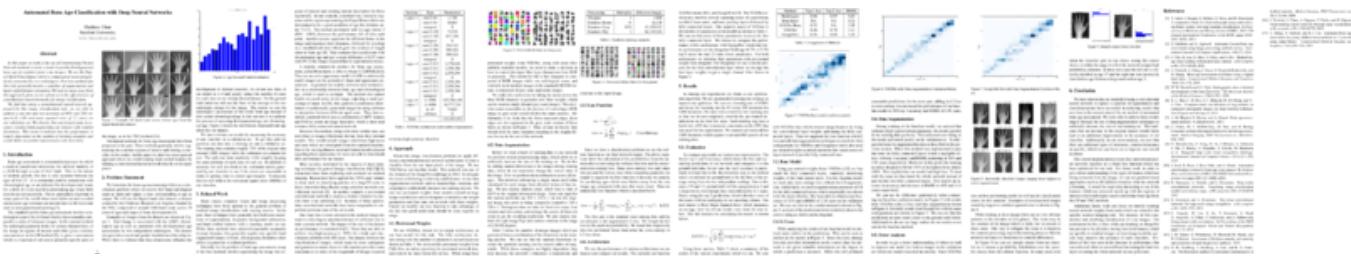
SherlockNet: Exploring 100 Years of Western Book Illustrations (reports2016/204_Report.pdf)

Brian Do, Luda Zhao, Karen Wang



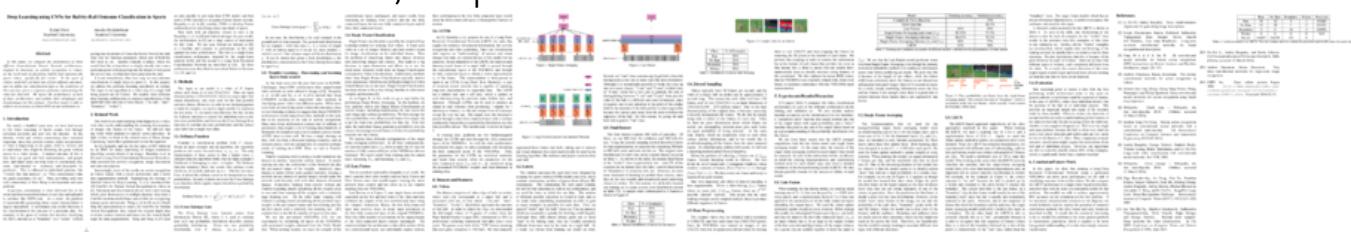
Automated Bone Age Assessment (reports2016/310_Report.pdf)

Matthew Chen



Ball-by-Ball Outcome Prediction for Cricket Matches (reports2016/273_Report.pdf)

Anusha Balakrishnan, Kalpit Dixit



Object Detection with Reduced Image Qualities (reports2016/287_Report.pdf)

Shiquan Wang, Hao Jiang



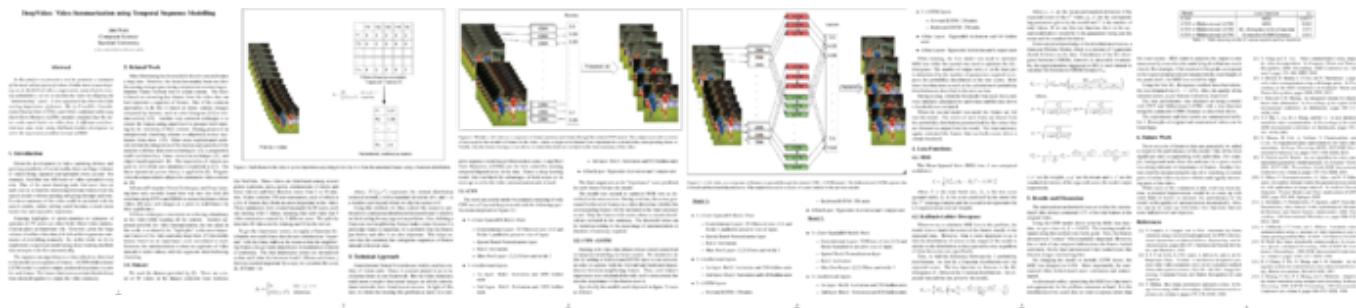
Pulmonary Nodule Classification with Convolutional Neural Networks (reports2016/324_Report.pdf)

Sheila Ramaswamy, Karen Truong



DeepVideo: Video Summarization using Temporal Sequence Modelling (reports2016/216_Report.pdf)

Juhi Amitkumar Naik



Deep Residual Learning for Direct Perception in Autonomous Driving (reports2016/123_Report.pdf)

Zuozhen Liu, Yixin Wang



Distilling Knowledge to Specialist ConvNets for Object Recognition with Clustered Classes (reports2016/120_Report.pdf)

Nathanael Romano, Robin Schucker



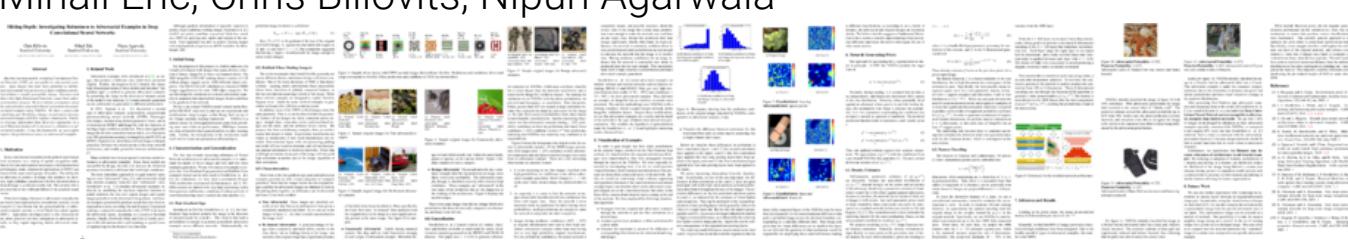
Measuring Cellular Semantic Distance with Siamese Neural Nets (reports2016/320_Report.pdf)

Morgan Paull



Hitting Depth: Investigating Robustness to Adversarial Examples in Deep Convolutional Neural Networks (reports2016/119_Report.pdf)

Mihail Eric, Chris Billovits, Nipun Agarwala



Using Convolutional and Recurrent Neural Networks on Magnetoencephalography Data to Assess Neuroplasticity (reports2016/334_Report.pdf)

Irán Román



Automatic Colorization with Deep Convolutional Generative Adversarial Networks (reports2016/224_Report.pdf)

Stephen Koo



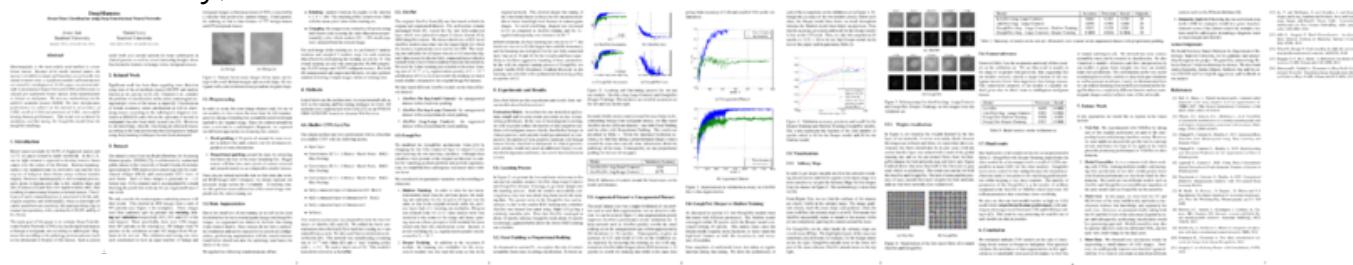
Disney Characters: Classifying Commercial Merchandise based on Cartoon Characters (reports2016/265_Report.pdf)

Paula Kusumaputri, Edwin Park



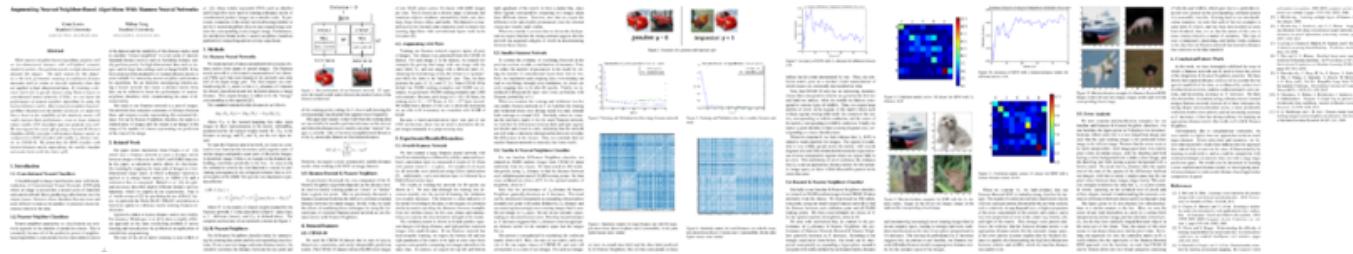
DeepMammo (reports2016/306_Report.pdf)

Daniel Levy, Arzav Jain



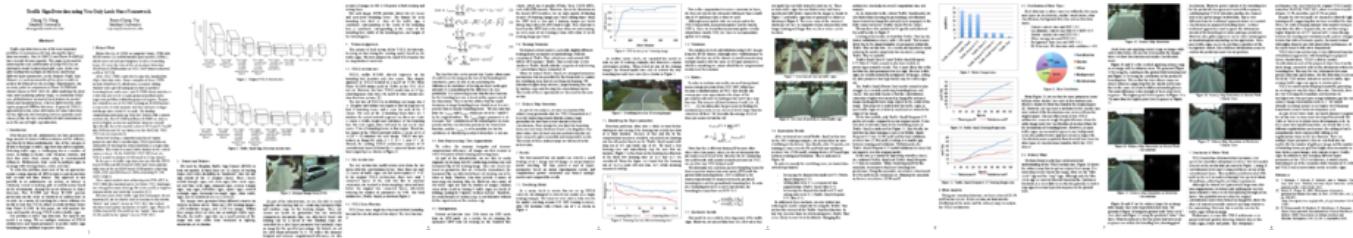
Augmenting Nearest Neighbor-Based Algorithms with Siamese Neural Networks (reports2016/258_Report.pdf)

Gene Lewis, Wilbur Yang



Traffic Sign Detection Using You Only Look Once Framework (reports2016/263_Report.pdf)

Royce Cheng-Yue, Chung Yu Wang



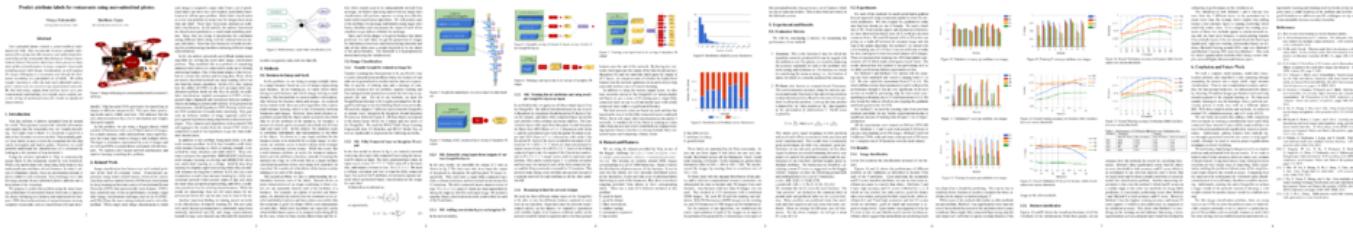
ResFuse and RefiNets: Enhanced CNN architectures for image classification (reports2016/277_Report.pdf)

Frederic Ren, Allen Nie, Nihit Desai



Predict attribute labels for restaurants using user-submitted photos (reports2016/025_Report.pdf)

Shubham Gupta, Vinaya Polamreddi



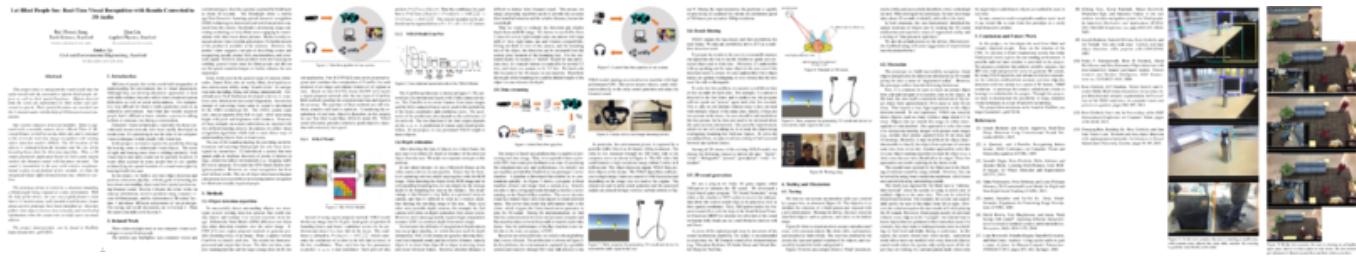
Using Convolutional Neural Networks to demystify aesthetic works of art (reports2016/210_Report.pdf)

Pujun Bhatnagar



Let Blind People See: Real-Time Visual Recognition with Results Converted to 3-D Audio (reports2016/218_Report.pdf)

Rui (Forest) Jiang, Qian Lin, Shuhui Qu



Lip-Reading with CNNs and LSTMs

(reports2016/217_Report.pdf)

Sameep Bagadia, Amit Garg, Jonathan Noyola



Tracking a Parking Lot's Empty Spaces Without Sensors

(reports2016/280_Report.pdf)

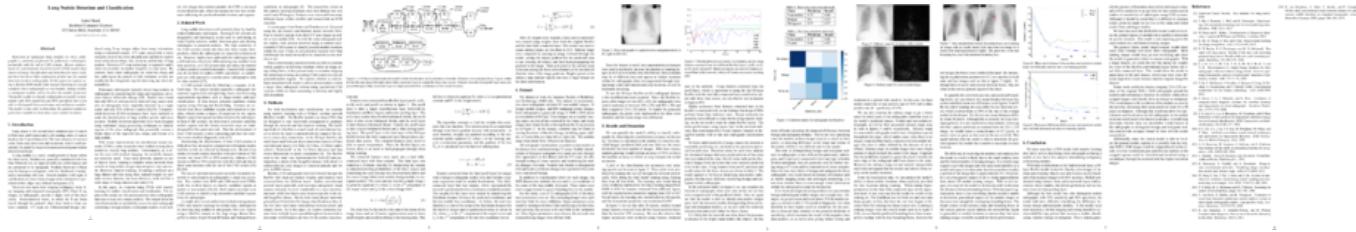
Jordan Cazamias, Martina Marek



Detecting and Classifying Lung Nodules in Chest

Radiographs (reports2016/313_Report.pdf)

Isabel Bush



Multi-Frame Video Super Resolution using Deep

Convolutional Networks (reports2016/212_Report.pdf)

Alex Greaves, Hanna Winter



Facial Affect Detection Using Convolutional Neural Networks (reports2016/011_Report.pdf)

Sherrie Wang



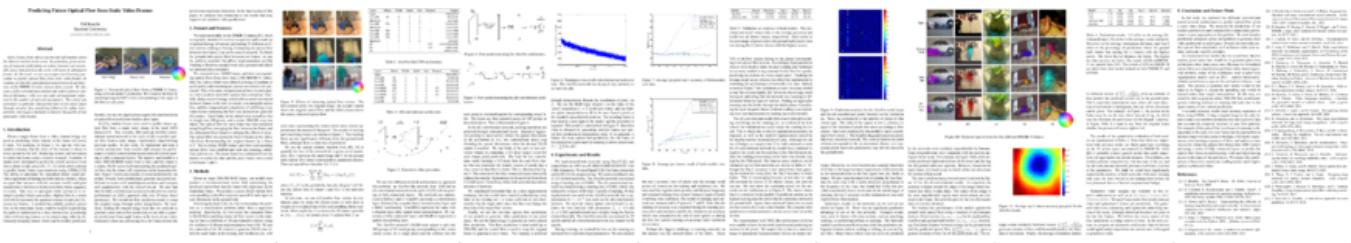
Solving human psychophysics tasks with 3D convolutional neural networks (reports2016/408_Report.pdf)

Daniel Birman, Dylan Cable



Predicting future optical flow from static video frames (reports2016/223_Report.pdf)

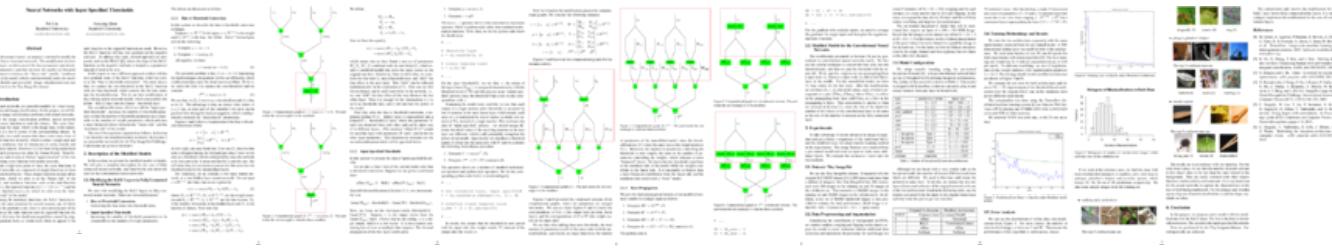
Pol Rosello



Neural Networks with Input Specified Thresholds

(reports2016/118_Report.pdf)

Fei Liu, Junyang Qian



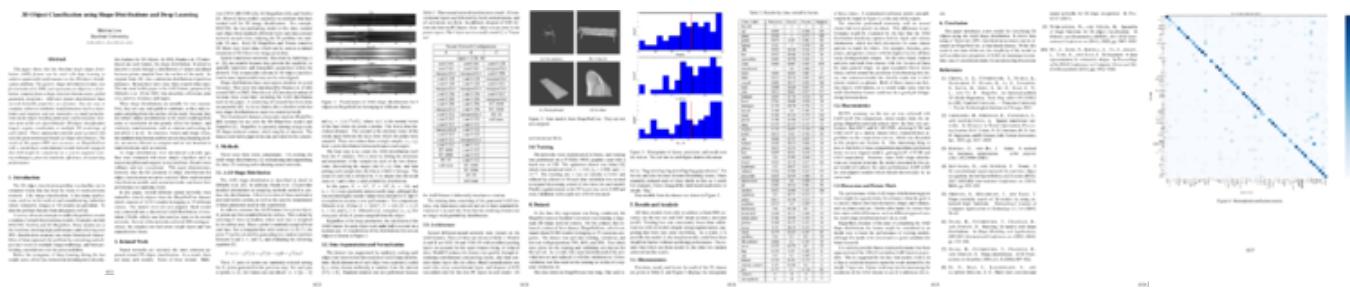
Recognizing Handwritten Japanese Characters Using Deep Convolutional Neural Networks (reports2016/262_Report.pdf)

Charlie Tsai



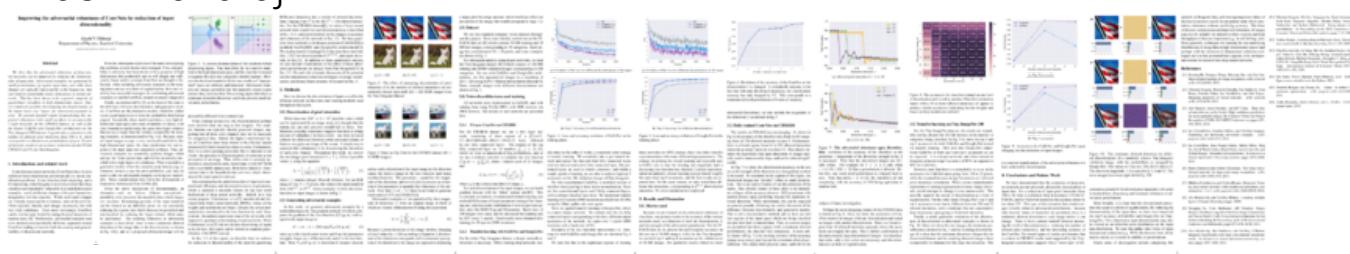
3D Object Classification using Shape Distributions and Deep Learning (reports2016/420_Report.pdf)

Melvin Low



Improving the adversarial robustness of Convnets (reports2016/103_Report.pdf)

Akash Maharaj



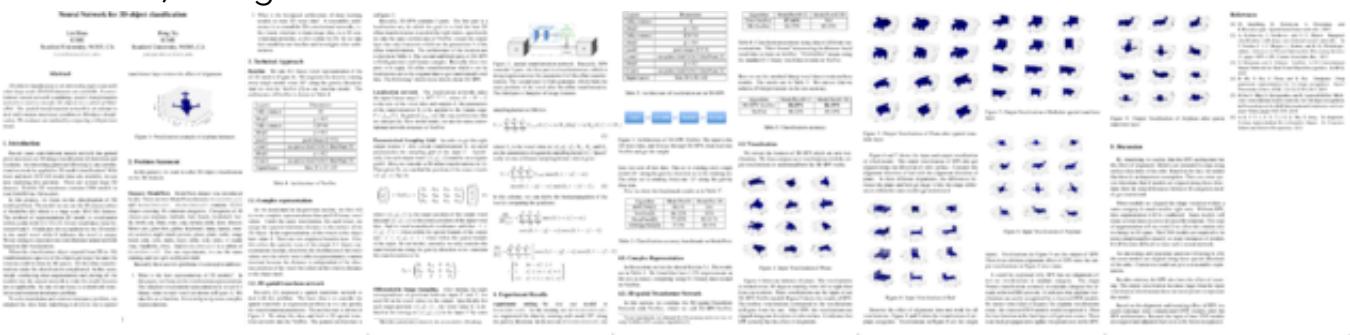
Playing Flappybird with Deep Reinforcement Learning (reports2016/111_Report.pdf)

Naveen Appiah Balaji, Sagar Vare



Neural Network for 3D object classification (reports2016/417_Report.pdf)

Lin Shao, Peng Xu



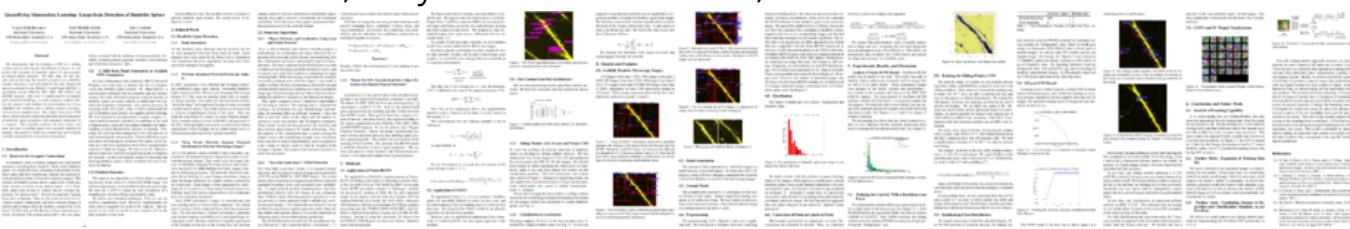
Facial Expression Recognition with Convolutional Neural Networks (reports2016/023_Report.pdf)

Vivek Choksi, Arushi Raghuvanshi



Quantifying Mammalian Learning: Large-Scale Detection of Dendritic Spines (reports2016/329_Report.pdf)

Seth Hildick-Smith, Ivaylo Bahtchevanov , John Lambert



Intelligent coloring of greyscale images using CNNs (reports2016/202_Report.pdf)

Tian Zhao, Ross Daly



Convolutional Neural Networks for Facial Age and Gender Prediction (reports2016/003_Report.pdf)

Ari Ekmekji



ambience_is_classy or good_for_kids? : Yelp Restaurant Classification Challenge (reports2016/015_Report.pdf)

Jee Ian Tam



Understanding Visual Art with CNNs (reports2016/201_Report.pdf)

Michael Baumer



Action Recognition and Detection Using CNNs and LSTMs (reports2016/221_Report.pdf)

Helen Jiang, Gary Thung



Facial keypoints detection using Neural Network (reports2016/007_Report.pdf)

Chenyue Meng, Shutong Zhang



Near Real-Time Super-Resolution Video (reports2016/203_Report.pdf)

David Zeng



Transforming How We Diagnose Heart Disease Using Convolutional Neural Networks (reports2016/327_Report.pdf)

Viswajith Venugopal, Swaroop Ramaswamy



Improving image feature representations with noisy similarity of text captions (reports2016/360_Report.pdf)

Dan Iter

A Recurrent Neural Game Engine (reports2016/101_Report.pdf)

David Gottlieb

Playing Tetris with Deep Reinforcement Learning (reports2016/121_Report.pdf)

Matt Stevens, Sabeek Pradhan

Quantifying Invariance in Convolutional Neural Networks (reports2016/107_Report.pdf)

Eric Kauderer-Abrams

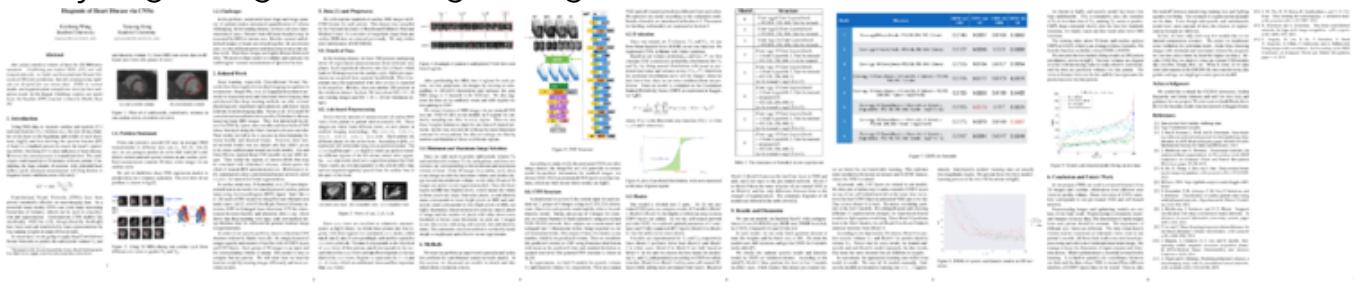
Yelp Restaurant Photo Classification (reports2016/013_Report.pdf)

Jade Huang



Diagnosis of Heart Disease via CNNs (reports2016/331_Report.pdf)

Yanyang Kong, Kaicheng Wang



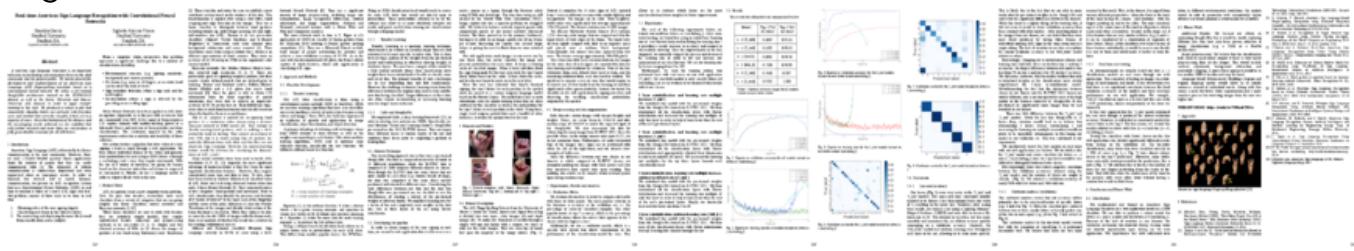
Single Image Depth Prediction Using Convolutional Neural Networks (reports2016/407_Report.pdf)

Ralph Ma, Jack Zhu



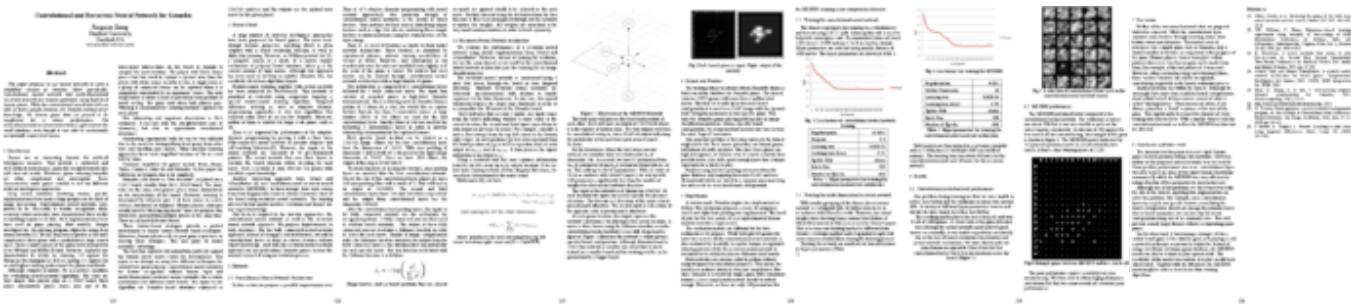
Read my Hands: Translating American Sign Language to Text (reports2016/214_Report.pdf)

Sigberto Alarcon Viesca, Brandon Garcia

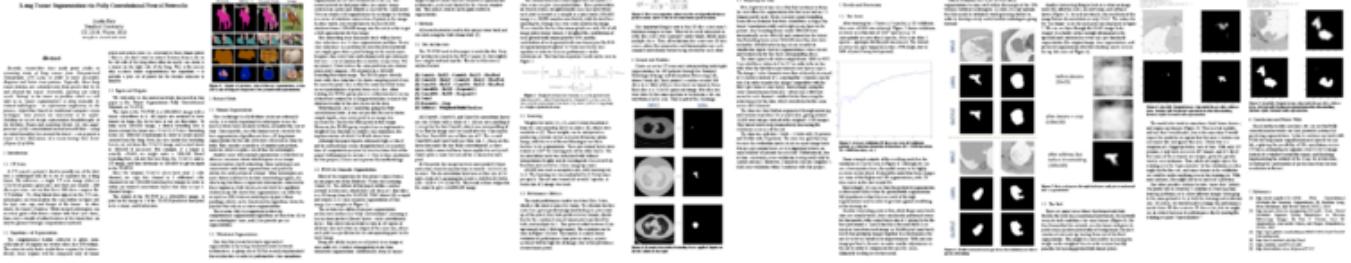


Game playing using recurrent neural networks (reports2016/109_Report.pdf)

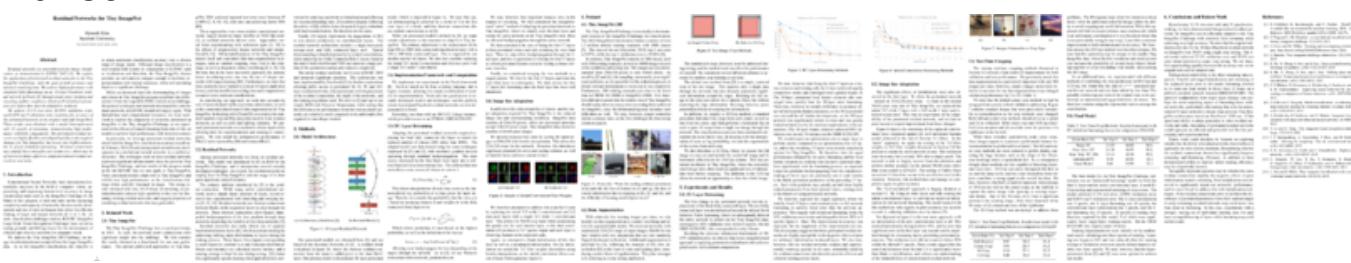
Nic Zhang



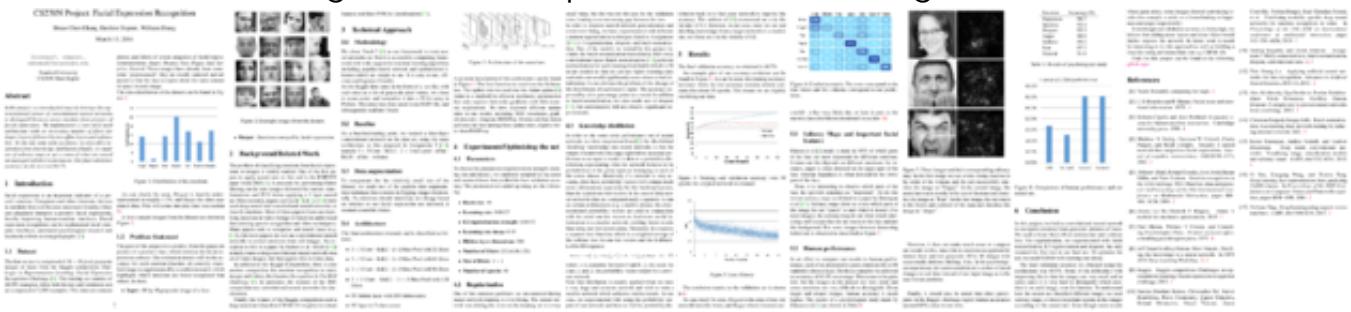
Fully Connected Convolutional Neural Networks for Lung Tumor Segmentation (reports2016/302_Report.pdf) Austin Ray



Residual Networks for Tiny ImageNet (reports2016/411_Report.pdf) Hansohl Kim



Facial Expression Recognition (reports2016/009_Report.pdf) Hsiao Chen Chang, Emilien Dupont, William Zhang



Geolocalization of Street View Images: From Pixels to Places (reports2016/422_Report.pdf)

Blake Wulfe, David Hershey



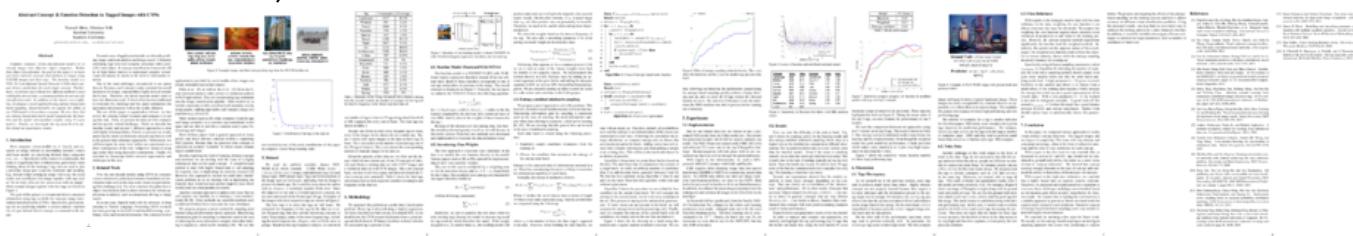
Classifying Restaurant Attributes With Deep Learning (reports2016/001_Report.pdf)

Pedro Garzon, Diveesh Singh



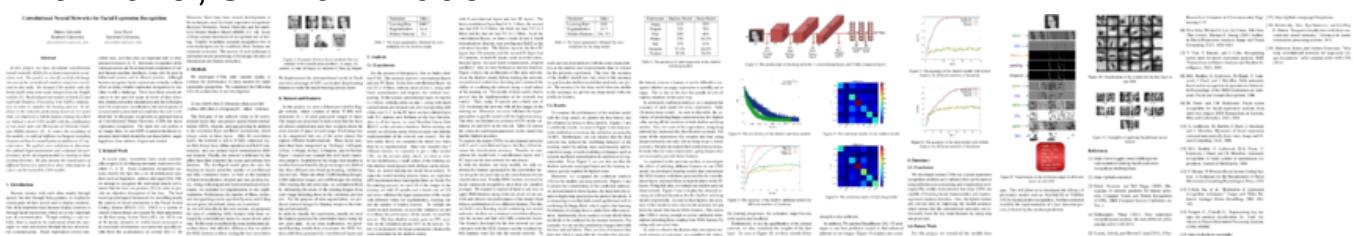
Abstract concept & emotion detection of tagged images with CNNs (reports2016/008_Report.pdf)

Youssef Ahres, Nikolaus Volk



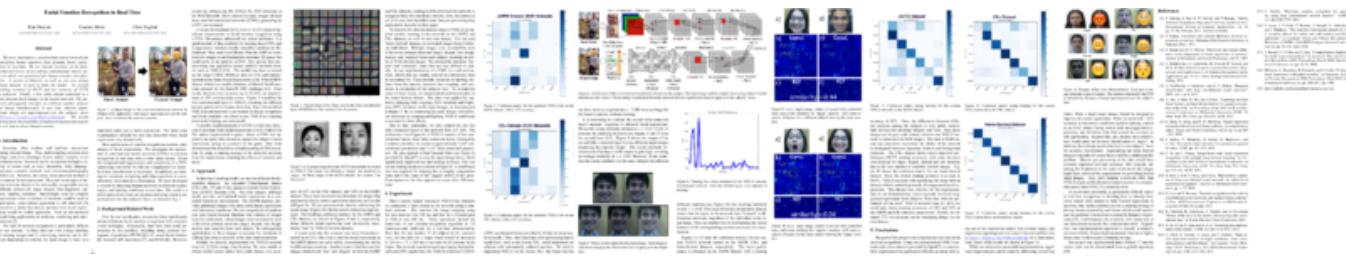
Convolutional Neural Networks for Facial Expression Recognition (reports2016/005_Report.pdf)

Azar Fazel, Shima Alizadeh



Facial Emotion Recognition in Real Time (reports2016/022_Report.pdf)

Dan Duncan, Chris English, Gautam Shine



A New Algorithm for Fully Automatic Brain Tumor Segmentation with 3-D Convolutional Neural Networks (reports2016/322_Report.pdf)

Christopher Elamri, Teun de Planque



Fast Unsupervised Object Localization (reports2016/285_Report.pdf)

Mihir Mongia, deepak menghani, Anjan Dwaraknath



Facial Keypoint Detection (reports2016/010_Report.pdf) Shayne Longpre, Ajay Sohmshetty



Cardiovascular Edge Detection using Neural Networks (reports2016/311_Report.pdf) Gabriel Maher

Convolution Neural Networks for Chinese Handwriting Recognition (reports2016/428_Report.pdf)

Xu Chen

Music Structure Segmentation with Convolutional Neural Networks (reports2016/220_Report.pdf)

Tim O'Brien

Facial Expression Recognition with Selected Partial Facial Feature Using CNN (reports2016/020_Report.pdf)

Yang Li, Chenjie Yang , Weier Wan

Brain tumour segmentation from MRI

(reports2016/328_Report.pdf)

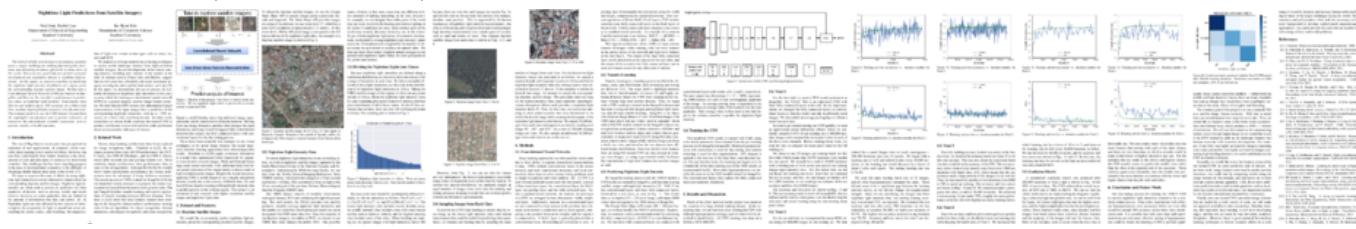
Raunaq Rewari



Nighttime Light Predictions from Satellite Imagery

(reports2016/423_Report.pdf)

Rachel Luo, Neal Jean, Jake Kim



Coloring Grayscale image (reports2016/205_Report.pdf)

Vivek Bagaria, Kedar Tatwawadi



Yelp Restaurant Photo Classification: MIML Approach with

CNNs (reports2016/017_Report.pdf)

Raghav Gupta, Pulkit Agrawal



Pedestrian Attribute Detection using CNN

(reports2016/255_Report.pdf)

Agrim Gupta, Jayanth Ramesh



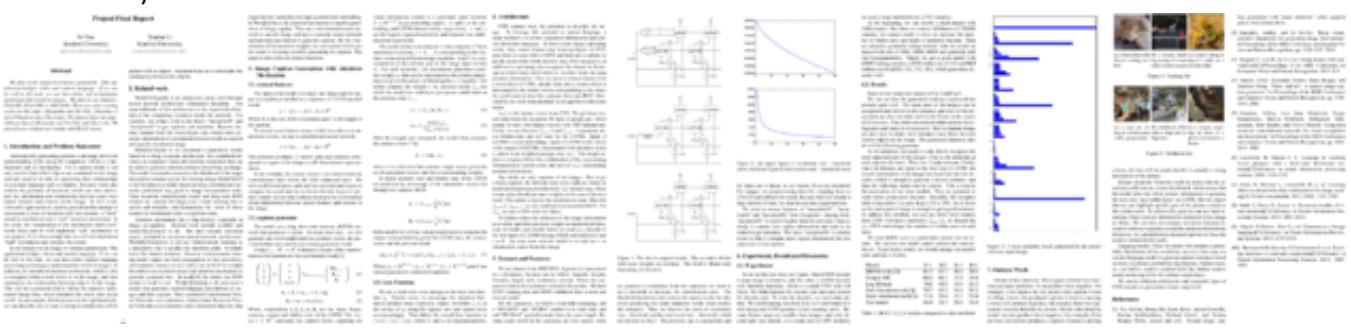
Multiple Object Recognition with Focusing and Blurring (reports2016/259_Report.pdf)

Holly Chiang, Yifan Ge, Connie Wu



Image Captioning with Attention (reports2016/364_Report.pdf)

Ye Tian, Tianlun Li



Machine Cashier (reports2016/260_Report.pdf)

Roy Chan



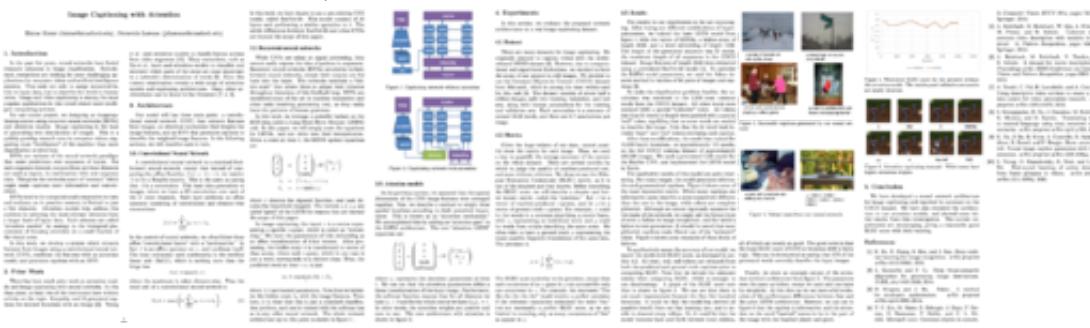
Tractable Neural Networks for Identity Recognition (reports2016/250_Report.pdf)

David Eng, Jenny Hong



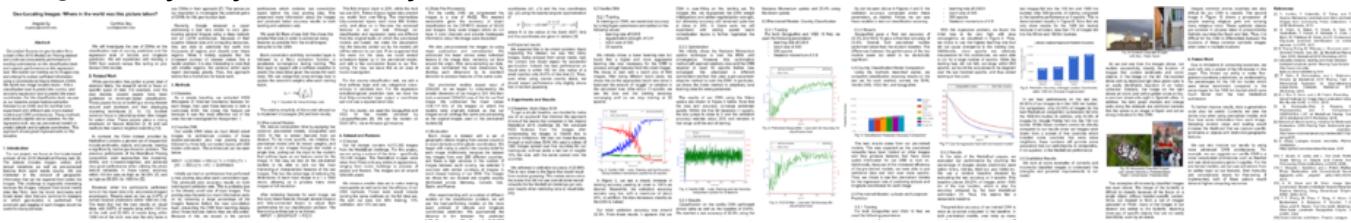
Models of Attention for Video Captioning (reports2016/362_Report.pdf)

Dieterich Lawson, Blaine Rister



Geolocating Images: Where in the world was this taken? (reports2016/404_Report.pdf)

Angela Sy, Cynthia Day



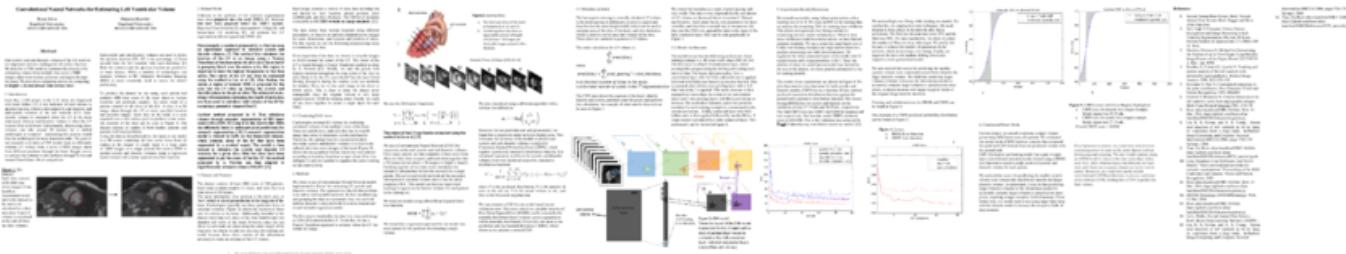
Flower Taxonomic Classification using CNNs (reports2016/253_Report.pdf)

Xavier Mignot, Maxwell Siegelman



Convolutional Neural Networks for Estimating Left Ventricular Volume (reports2016/323_Report.pdf)

Ryan Silva, Maksim Korolev



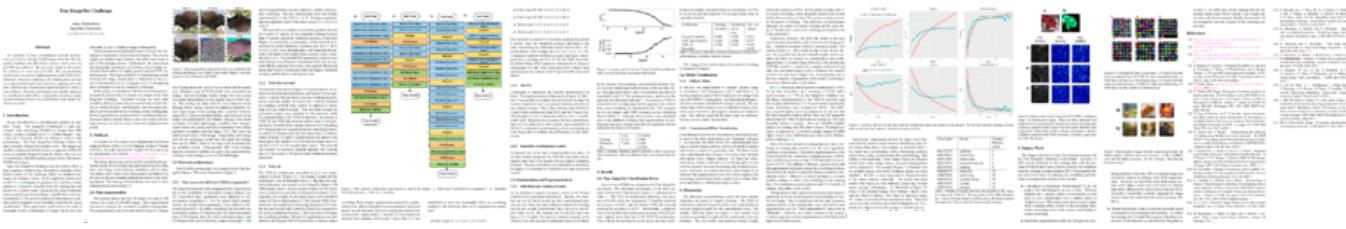
Going Deeper on the Tiny Imagenet Challenge (reports2016/405_Report.pdf)

Andrew Zhai



Tiny ImageNet Visual Recognition Challenge (reports2016/401_Report.pdf)

Anna Shcherbina



Convolutional Neural Network for Traditional Chinese Calligraphy Recognition (reports2016/257_Report.pdf)

Boqi Li



Image Detection Techniques on the Daimler Monocular Pedestrian Detection Benchmark Data (reports2016/288_Report.pdf)

Christopher Ling

Study of Residual Networks for Image Recognition (reports2016/264_Report.pdf)

Mohammad Sadegh Ebrahimi, Hossein Karkeh Abadi

Using Convolutional Neural Network for the Tiny ImageNet Challenge (reports2016/425_Report.pdf)

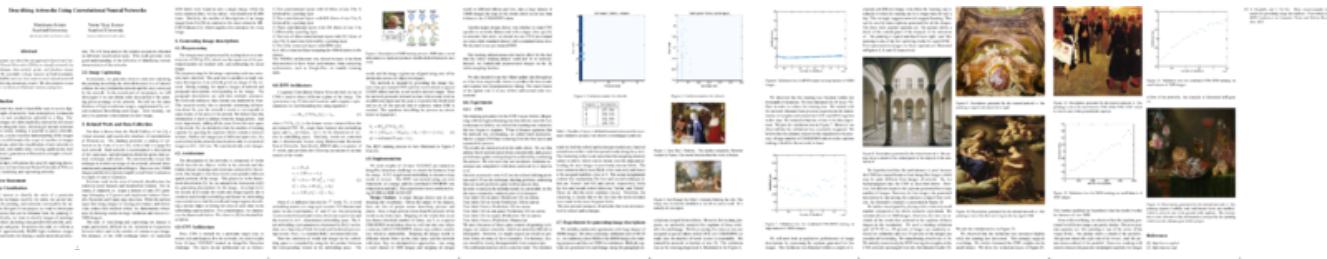
Jason Ting

Neurofusion: Fusing MEG and EEG Data (reports2016/321_Report.pdf)

Paul Warren

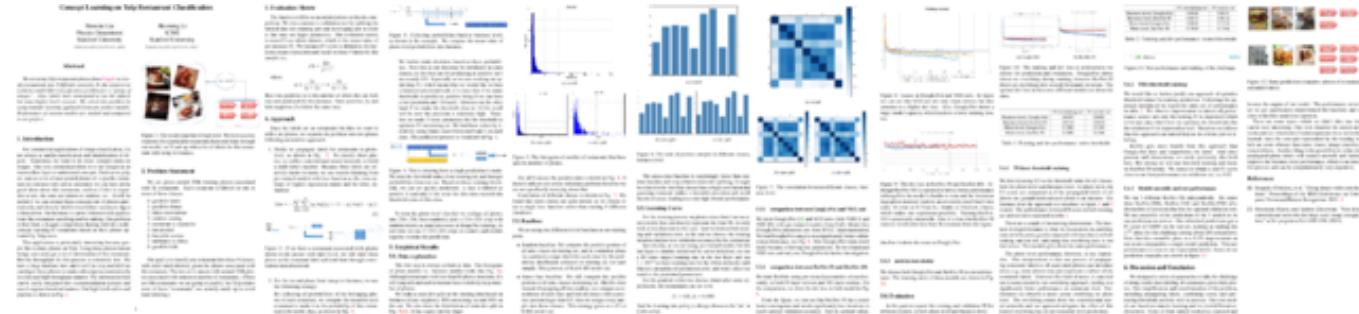
Virtual art museum (reports2016/200_Report.pdf)

Manikanta Kotaru, Varun Vijay Kumar



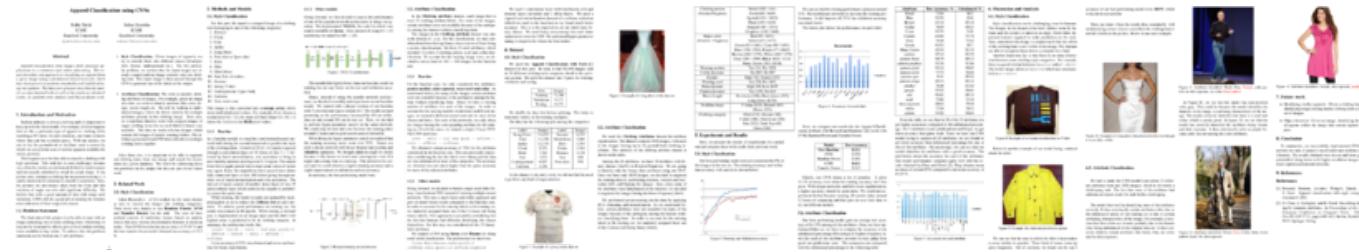
Concept Learning on Yelp Restaurant Classification (reports2016/024_Report.pdf)

Shenxiu Liu, Haoming Li



Apparel classification using CNNs (reports2016/286_Report.pdf)

Rohit Patki, Suhas Suresha



A network combination of ResNet and GoogLeNet to tackle Tiny ImageNet (reports2016/410_Report.pdf)

William Shen



Yelp Restaurant Photo Classification Challenge (reports2016/014_Report.pdf)

Rajarshi Roy

Application of Convolved Neural Networks for Pedestrian Detection (reports2016/254_Report.pdf)

Anil Variyar

Classifying Restaurants with Yelp Photos (reports2016/021_Report.pdf)

Mike Yu, Emma Marriott, Aaron Zweig

Tiny ImageNet Challenge (reports2016/419_Report.pdf)

Vani Khosla

Identifying the Higgs Boson (reports2016/300_Report.pdf)

Anton Apostolatos, Leonard Bronner



Logo Recognition in Real World Images

(reports2016/268_Report.pdf)

Naveen Arivazhagan

Logo Recognition

Manav Arivazhagan
Department of Computer Science
Stanford University

Number	Name	Report ID	Report URL
1	Logo Recognition	268	http://cs231n.stanford.edu/reports2016/268_Report.pdf
Total		1	

1. Introduction
In this project we work on logo recognition in real world images. We have used a convolutional neural network to identify the logo in the image. The network takes the image as input and outputs the probability of each logo being present in the image.

2. Related Work
There are many different approaches to logo recognition. One approach is to use a template matching algorithm. This algorithm compares the logo with every possible location in the image and finds the best match. Another approach is to use a neural network. This network takes the image as input and outputs the probability of each logo being present in the image.

3. Problem Setup
The problem setup for this project is as follows. We have a dataset of images containing logos. The dataset has been split into training and testing sets. The training set is used to train the neural network. The testing set is used to evaluate the performance of the network.

4. Performance Metric
The performance metric for this project is accuracy. Accuracy is defined as the percentage of correctly identified logos in the testing set.

5. Results
The results of the project are as follows. The neural network achieves an accuracy of approximately 90% on the testing set. This is a good result, but there is still room for improvement.

6. Evaluation
The evaluation of the project is as follows. The neural network is evaluated on a variety of datasets. The results show that the network performs well on most datasets, but there are some datasets where the performance is not as good.

7. Future Work
There are several areas for future work. One area is to improve the performance of the neural network. Another area is to extend the project to recognize more types of logos. A third area is to use the neural network for other applications, such as image segmentation.